

THE MISSION GREY GUILD NEWSLETTER

DOUBLE ISSUE, 2026



HEALTH OF NATIONS

THE HIDDEN ECONOMICS
OF HUMAN CAPACITY

Health of Nations

The hidden economics of human capacity

PEOPLE
CAPACITY
PROSPERITY
A STRONGER
TOMORROW

Healthier people.
Stronger economies.
Brighter futures.

TABLE OF CONTENTS

#	ARTICLE / SECTION	PAGES	KEY THEMES	KEY SECTORS (2-3)	KEY AUDIENCE (2-3)
1	 FOREWORD: The Nation as Patient	4 – 7	- Human capacity and national resilience - Limits of economic indicators - Hidden reserves and fragility - Productivity vs. long-term capacity	 Public sector  Healthcare  Industry	 Policymakers Business leaders Strategic planners
2	 The GLP-1 Economy Is Escaping Healthcare... And Rewriting Consumer Markets	8 – 9	- GLP-1 drugs and wider economic impact - Changing consumer behaviour - Health, productivity and spending - Ripple effects beyond healthcare	 Healthcare & pharma  Consumer goods & retail  Financial services	 Investors Marketing & strategy Business leaders
3	 The Workforce That Shows Up	10 – 19	- Labour force participation and reliability - Health, wellbeing and attendance - Demographic pressures - What enables people to keep working	 Manufacturing  Healthcare  Professional services	 HR leaders Operations Business executives
4	 Five Hidden Indicators of Fatigue	20	- Early warning signs in human capacity - Measuring stress, strain and fragility - Indicators beyond GDP - Practical signals for organizations	 All sectors (cross-industry)	 HR and people analytics Executive leadership
5	 Reading Human Capacity	21 – 28	- Measuring and analysing human capacity - Data, metrics and new tools - Linking health, skills and economic output - From measurement to action	 Public sector  Industry  Financial services	 Analysts & researchers Strategy teams Policymakers
6	 The Vacancy That Refuses to Die	29 – 30	- Persistent labour shortages - Skills mismatches - Demographic change and ageing - Structural vs. cyclical factors	 Manufacturing  Technology  Education	 HR and talent acquisition Operations Business leaders
7	 The Missing Apprentice	31 – 36	- Impact of automation and AI on early careers - Loss of learning opportunities - Building the next generation of expertise - Risks to long-term institutional knowledge	 Technology  Manufacturing  Education	 HR and learning & development Educators Business leaders
8	 Scenario Piece: Four News Dispatches from the Future (2035)	37 – 41	- Alternative futures for human capacity - Consequences of inaction - Opportunities for renewal - Strategic choices today	 Public sector  Industry  Technology	 Business leaders Strategic planners Policymakers

► THIS WEEK IN BUSINESS RISK

Serious risks. Unexpected questions.
No crystal ball required.



A more connected world means more complex risks.

From supply chains and critical minerals to AI, energy, trade and geopolitics — the forces shaping tomorrow's business environment are moving faster, intersecting in new ways and hitting closer to home.

This Week in Business Risk, the Mission Grey Guild's podcast hosted by **Rwizi Rweizooba Ainomugisha**, cuts through the noise to help you understand what really matters — and what to watch next.



A war can become a shipping problem.



A technology rivalry can become a supply-chain problem.



A mining law can change the strategic outlook for an entire industry.



A seemingly obscure geopolitical development can suddenly matter to a company thousands of kilometres away.

RECENT EPISODES, REAL CONVERSATIONS

1



Mission Grey CEO Lauri Tähtinen discusses external intelligence, human-machine teaming and the changing nature of corporate decision-making — and what it means for business in an era of **great-power competition**.

2



Mercy Kuo, Mission Grey's Head of Strategy and a contributing author at *The Diplomat*, discusses the US-China AI race, what it means for global supply chains, and how companies can prepare for **a more fragmented world**.

3



Critical minerals, India's mining reforms, emerging India-Africa trade links and the geopolitical battle around AI — and what they mean for companies, investors and **entire industries**.



LISTEN TO
**THIS WEEK IN
BUSINESS RISK**

ON YOUTUBE, SPOTIFY,
APPLE PODCASTS AND
OTHER MAJOR PODCAST
PLATFORMS.



Rwizi brings experts into the conversation to work through those connections — with different perspectives, practical insights and a few uncomfortable questions.



The result is not another podcast in which two people spend an hour agreeing that the world is becoming increasingly uncertain. Episodes move quickly from the big strategic question to the practical one: **so what should a decision-maker actually be watching?**



That combination is what makes *This Week in Business Risk* different.



It is serious without being solemn. Geopolitical without being trapped in geopolitics. And business-focused without pretending that business exists separately from the rest of the world.



The Guild's experts bring the knowledge. Rwizi brings the questions.



THE REST IS CONVERSATION.



YouTube




Spotify



Apple Podcasts

+ MORE



FOREWORD

The Nation as Patient

“The first wealth is health.”

— Ralph Waldo Emerson

We have become very good at measuring the temperature of economies. Every month, sometimes every hour, we consult the instruments. Inflation is up or down. Growth has been revised by a decimal point. Bond yields move. Purchasing managers report optimism or fatigue. Unemployment remains unexpectedly low, or begins to rise. Oil prices react to war. Currencies respond to uncertainty. Governments publish new forecasts, central banks adjust their language, and economists gather around the charts as though around a bedside monitor.

This is not foolishness. Temperature matters. A fever can kill, and so can a collapse in circulation. The macroeconomic indicators we have developed over decades remain indispensable. GDP, inflation, employment, productivity, debt, investment, and trade tell us things we need to know.

But they do not tell us everything we need to know.

A country can have respectable growth and declining capacity. It can have high employment while its hospitals struggle to staff wards, schools struggle to keep teachers in classrooms, manufacturers struggle to replace retiring technicians, and public administrations struggle to retain institutional

memory. It can report longer life expectancy while more of those additional years are lived with illness, disability, or dependence. It can invest heavily in artificial intelligence while weakening the human pipelines through which skills, judgment and leadership are traditionally acquired.

The central question, therefore, is not simply whether an economy is growing. It is whether a country can still do the work it needs to do.

That question lies at the heart of this double issue of the Mission Grey Guild.

We call it ***Health of Nations: The hidden economics of human capacity*** because the language of health helps us examine something that conventional economic statistics often obscure. Nations, like people, can survive for a surprisingly long time while consuming their reserves. They can compensate for weaknesses. They can borrow against the future. They can function through exhaustion. They can maintain an appearance of stability even as the margin for error becomes progressively thinner.

The reserves of a nation are not only financial. They include accumulated skills, functioning institutions, social trust, demographic momentum, experienced workers, healthy working years, infrastructure, professional cultures, and the ability to pass knowledge from one generation to the next. These reserves can be consumed quietly.

A hospital that relies permanently on overtime is consuming a reserve. A school system that depends on exhausted teachers is consuming a reserve. An industry that loses experienced technicians faster than it can train replacements is consuming a reserve. A public administration that repeatedly loses experienced officials without transferring their knowledge is consuming a reserve. A society in which middle aged workers simultaneously carry the demands of employment, children and ageing parents is consuming a reserve. None of these developments necessarily appears immediately in GDP.

That is precisely the problem.

For decades, efficiency was often treated as synonymous with strength. We optimized supply chains, reduced inventories, outsourced functions, compressed staffing and sought maximum output from minimum inputs. Much of this produced enormous gains. Yet the experience of recent years has reminded us that efficiency without reserves can become fragility.

The same principle applies to human capacity. A workforce can be statistically employed and operationally brittle. A country can have fewer unemployed people while simultaneously having fewer people capable of filling the jobs that matter most. A company can have access to advanced technology while losing the junior employees who would once have learned the profession by performing the repetitive tasks that technology now performs.

This creates an uncomfortable paradox: technology may increase productivity while weakening the replenishment of human capability.

The problem is not artificial intelligence itself. The problem is what happens when automation removes not only undesirable work, but also the early career work through which people learn how a profession actually functions. In law, finance, consulting, administration, and many other fields, entry level tasks have traditionally been the first rung of a much longer ladder. If that rung disappears, the consequences may not be visible immediately. Five or ten years later, however, an

organization may discover that it has plenty of senior expertise but far fewer people capable of becoming the next generation of experts.

A nation can therefore become more efficient today while becoming less capable tomorrow. This is why the familiar language of resilience deserves closer examination.

A nation is not resilient merely because it has not yet failed. Resilience means having enough reserve capacity to absorb disruption without exhausting the people and institutions on which the system depends. It means being able to replace skills without losing continuity. It means extending longevity without simply extending years of frailty. It means digitizing without systematically deskilling. It means supporting care without quietly transferring its costs to households. It means maintaining enough healthy working years, trained people, and social trust to keep essential systems functioning when circumstances become difficult.

The demographic dimension is particularly important. Ageing is often discussed as though it were itself a diagnosis. It is not. A society with fewer young people is not automatically condemned to decline. Scarcer labour can create incentives for investment, automation, and productivity growth. Longer lives can represent one of the great achievements of modern civilization.

But demographic change changes the conditions under which everything else operates. When fewer workers support more dependants, when experienced employees retire faster than replacements can be trained, and when care responsibilities increasingly fall on working age households, the pressure appears throughout the economy. Labour markets tighten. Public expenditure rises. Employers compete more aggressively for scarce skills. Families absorb costs that do not always appear in national accounts. Immigration becomes not merely a political question, but an economic necessity in many countries.

The result is a new competition for human capacity. Businesses compete for workers. States compete for taxpayers. Health systems compete for nurses and doctors. Universities compete for researchers. Cities compete for skilled migrants. Defence organizations compete for people capable of operating increasingly sophisticated systems.

The countries that understand this competition will not necessarily be those with the largest populations. They will be those capable of making the best use of the people they have, attracting the people they need, and preserving the health and skills of both.

For business leaders, this is not a peripheral social issue. It is a strategic variable.

The health of a nation's population affects labour supply, productivity, consumption, insurance costs, taxation, public finances, service delivery, and the state's ability to execute policy. It affects whether a factory can operate at full capacity, whether a hospital can maintain safe staffing, whether a municipality can deliver basic services, and whether an armed force can sustain readiness.

At some point, every economic strategy encounters the same stubborn question:

Who, exactly, will do the work?

Money cannot answer that question by itself. Nor can technology.

Artificial intelligence may transform productivity. Automation may compensate for labour shortages. Migration may replenish populations. Better healthcare may extend healthy working lives. Education may increase human capital. None of these solutions is sufficient in isolation.

The real challenge is to understand how they interact.

That is what this double issue attempts to do.

We are not approaching health primarily as a healthcare question, nor ageing merely as a demographic question. We are examining a harder political economy of human capacity. Where are labour shortages actually shortages of people? Where are they shortages of skills? Where are they shortages of healthy years? Where are they symptoms of institutions that have become too demanding of the people on whom they depend?

And perhaps most importantly, are governments and companies managing these pressures, or merely becoming increasingly sophisticated at accounting for them?

Good diagnosis begins by refusing comforting proxies.

GDP is not capacity. Employment is not readiness. Longevity is not vitality. AI investment is not human renewal. Healthcare spending is not the same thing as a healthy society.

The question beneath all of them is simpler, sterner and more consequential. When a nation needs people to build, teach, care, maintain, govern, invent and defend, are enough healthy, skilled and available people still there? That is the intellectual spine of this issue. Everything else follows from it.

Pekka Virkki

Guildmaster & Editor-in-Chief, Mission Grey Guild



The GLP-1 Economy Is Escaping Healthcare... And Rewriting Consumer Markets

One of the most important business stories of late 2026 is that GLP-1 weight-loss drugs are no longer just a pharmaceutical trend. They are becoming an economy in their own right, with implications reaching into retail, food, insurance, supplements, wellness, and employer benefits.

The immediate trigger is obvious: access has broadened, even if unevenly. In the United States, eligible Medicare beneficiaries can now obtain certain GLP-1 medicines for about \$50 a month through the Medicare GLP-1 Bridge, compared with prior costs above \$1,000 a month. At the same time, U.S. prescription drug prices were reported down 3.1% year on year in July, the sharpest annual drop in more than six decades, according to CMS leadership.

Yet this is not a story of universal access. It is a story of expansion colliding with affordability pressure. California is considering state-backed action to lower GLP-1 prices after prescriptions for weight loss rose from 20,000 to 700,000 between 2018 and 2023 across Medi-Cal and commercial plans. Several U.S. states have already cut Medicaid coverage for obesity use because of budget strain. Employers are also becoming more selective: the share covering GLP-1 drugs for weight loss fell from 72% in 2025 to 60% in 2026, while 14% already have or plan to drop coverage in 2027.


This matters because employer healthcare costs are projected to rise 9.2% in 2027 without further cost controls, with pharmacy spending expected to increase 12%.


What makes this trend strategically important is its spillover into mainstream consumer markets. A secondary ecosystem is forming around the drugs: nutrition products, post-treatment support, telehealth, diagnostics, coaching, and specialized retail. In the UK, Holland & Barrett has launched FORMI Weight Loss Complex for consumers seeking non-pharmaceutical support or nutritional support after GLP-1 treatment, while also partnering with online pharmacy Phlo on clinical weight-management services. One estimate puts the global GLP-1 support supplement market at roughly \$1.8 billion in 2025, rising to \$8.9 billion by 2034, implying around 18.5% compound annual growth.

For business, the real significance is second-order demand change. Food and beverage companies may face pressure in categories linked to snacking, indulgence, and oversized portions. Apparel retailers may need to rethink sizing mixes and replenishment patterns. Wellness, beauty, fitness, and preventive health brands may benefit as spending shifts toward self-optimization and medically adjacent consumption. Insurers, pharmacy benefit managers, and self-insured employers face the most immediate cost challenge, but consumer-facing sectors should not assume they are insulated.

The most prudent interpretation is neither hype nor dismissal. A high-adoption GLP-1 scenario remains **Likely** in affluent and insured segments, especially in North America and parts of Europe. But a coverage backlash is also **Possible**, as payers push back and employers narrow benefits. That means companies should plan for an uneven world: sustained growth in medically enabled weight management, but patchy access, fluctuating reimbursement, and strong demand for non-drug support products.

In short, GLP-1s are becoming a cross-sector business variable. The winners will not only be drugmakers. They will also include the companies that understand how a health innovation changes consumer behavior, pricing power, product design, and benefit economics far beyond healthcare itself.





The Workforce That Shows Up

For much of the industrial era, productivity improvement meant producing more with the same number of people.

The next frontier may be different. It may mean getting more sustainable productive capacity from the people who are already employed.

That does not mean squeezing more out of them. It means reducing the amount of capacity lost to preventable illness, poor job design, unnecessary physical strain, bad management, untreated mental-health problems and rigid working arrangements.

The distinction is crucial. The old productivity question was:

How much can a worker produce?

The emerging question is:

How much productive capacity can the organisation sustain without consuming the worker who provides it?

That is a more difficult question. It is also a more strategic one.

The demographic arithmetic is increasingly unforgiving. When populations age and labour markets tighten, companies cannot assume that another worker will always be available to replace the one who burns out. They cannot assume that a vacancy will be filled next month. They cannot assume that an experienced employee who leaves can easily be replaced with a younger one. And they cannot assume that a low absence rate means that their workforce is healthy. The scarce resource may increasingly be **healthy attention, reliable judgement, physical resilience, and sustainable human capacity**.

Presenteeism, hidden incapacity, and the productivity problem nobody measures properly

There is a peculiar success story in modern labour markets.

People have jobs. Unemployment is low. Companies complain about hiring. Governments worry about labour shortages. Yet some of the same economies that cannot find enough workers are quietly losing working capacity among the people who are already employed.

They are not necessarily absent.

They are at their desks. On the factory floor. In the classroom. Behind the hospital counter. Driving the bus. Answering emails. Joining the meeting. Opening the shop.

They are simply not fully there.

This is the territory of *presenteeism*, the phenomenon in which people attend work despite being ill or impaired, with a resulting reduction in their ability to perform effectively. It is the mirror image of *absenteeism*, but a much harder one to see.

Absence leaves a hole in the schedule. Presenteeism leaves a person in the schedule while quietly reducing the capacity of the system around them.

That distinction is becoming economically important.

The World Health Organization estimates that depression and anxiety alone result in around 12 billion lost working days globally each year, with an annual productivity cost approaching US\$1 trillion. It has been evaluated that the productivity cost of presenteeism may be around 1.5 times that of sickness absence. That figure should not be interpreted as a universal accounting constant, but it illustrates the central problem: what organisations can easily count may not be what costs them the most.

The workforce problem, in other words, may not simply be that too few people are available. It may also be that too many of the people who are available are working below sustainable capacity. And unlike an empty vacancy, that loss rarely appears on a dashboard.

The worker who is technically present

The conventional labour-market model is remarkably binary: a person is either employed or unemployed. At a company level, an employee is either at work or absent. A sickness certificate turns a person into an absence statistic. A resignation turns them into a vacancy.

But human capacity does not work that way. Someone can have a migraine and still answer emails. Someone can have chronic back pain and still operate a machine. Someone experiencing depression can still attend a meeting. A nurse can come to work exhausted. A teacher can enter a classroom after sleeping badly for weeks. A manager can remain permanently available while becoming progressively worse at making decisions.

None of these people necessarily qualifies as absent, yet their productive capacity has changed.

The distinction matters because modern economies increasingly depend on forms of work in which the quality of cognition, attention, judgement, communication, and emotional regulation matters as much as physical presence.

A worker who is 20% less attentive is not necessarily producing 20% less output in any easily measurable sense. The cost may instead appear as a missed opportunity, a bad decision, a slower customer interaction, an avoidable error, a delayed project, or an additional burden transferred to colleagues. In a factory, the consequences can be visible. In a knowledge organisation, they can remain invisible for months.

This is one reason presenteeism is so difficult to manage. It does not announce itself as a productivity problem. It often looks like commitment.

The employee came to work. The meeting happened. The inbox was cleared. The shift was covered. The spreadsheet was submitted. The customer was served. Everything appears normal... Until it is not.

Organisations naturally prefer measurable problems.

Absenteeism can be counted, but presenteeism is different. The employee is still technically supplying labour. The organisation therefore continues to count the person as part of its productive capacity, even when illness, exhaustion, or psychological strain has reduced what that person can actually contribute. This creates a particularly dangerous accounting illusion: the payroll says 100%, but the actual capacity may be considerably lower.

The important point is that this invisible category exists. And once it exists, the traditional productivity calculation begins to look incomplete. A company may celebrate a low sickness-absence rate while having an exhausted workforce. It may congratulate itself on attendance while losing productivity through errors, disengagement, slower execution, and staff turnover.

A country may have a high employment rate while simultaneously accumulating a population of people who are employed but increasingly unable to sustain the work they are doing.

That is not a contradiction. It is a measurement problem.

The economics of showing up sick

Presenteeism is sometimes treated as an individual behavioural problem. Why did the employee come to work? Perhaps they felt responsible. Perhaps they feared falling behind. Perhaps they were worried about their job. Perhaps the workplace culture rewards endurance. Perhaps they did not want colleagues to carry their workload. Perhaps they simply could not afford to stay home.

But the research suggests that presenteeism has multiple causes. A major meta-analysis covering 109 samples and almost 176,000 people identified a broad set of individual, occupational, and organisational factors associated with going to work while ill.

This matters because a company can inadvertently manufacture presenteeism.

Consider a department already operating with too few people. One employee becomes ill. The employee knows that staying home will increase pressure on colleagues. The manager knows that replacing the employee is difficult. The employee therefore comes to work. The organisation gets through the day. But the underlying problem has not disappeared.

The employee may recover more slowly. Performance deteriorates. Another colleague absorbs the additional workload. Stress increases. The probability of another absence rises. The system has transformed one absence into a distributed productivity loss. This is particularly relevant in sectors already suffering from labour shortages.

The irony is uncomfortable: the fewer people an organisation has, the greater the temptation to make the people it does have work through their reduced capacity. That may preserve the schedule today while damaging the workforce tomorrow.

When the labour shortage creates the health shortage

The relationship between labour scarcity and ill health can become circular.

A hospital cannot recruit enough nurses. Existing nurses work harder. Workloads rise. Recovery time falls. Stress accumulates. Sickness absence increases. The remaining workforce becomes even more stretched. More employees begin working while unwell because there is nobody to replace them. Productivity falls.

The organisation still needs more labour, but the original shortage becomes harder to solve. The same logic can apply in education, social care, transport, hospitality, manufacturing, and public administration.

This is why sickness absence should not always be interpreted as evidence that workers are withdrawing from the labour market; sometimes it is evidence that the labour market is consuming too much capacity from too few people.

Germany provides a useful contemporary example. Average sickness absence reached about 22 days per worker in 2024, according to data cited in a recent *Reuters* report. The burden was particularly pronounced in health and social work and education. Mental-health-related cases represented a relatively small share of cases but generated much longer absences, while days lost because of mental illness had risen sharply since 2014. Researchers have argued that understaffing, emotional strain, and working conditions are part of the explanation, rather than employee behaviour alone.

That distinction is strategically important. If the problem is interpreted as excessive absence, the obvious response is stricter absence control. If the problem is interpreted as accumulated system strain, the response looks different.

The first approach asks:

How do we get people back to work faster?

The second asks:

Why are so many people becoming unable to sustain work in the first place?

Those are not the same question.

Spain shows the scale of the problem

Spain illustrates how increasing absence from work can become a macroeconomic issue.

According to *Randstad Research*, overall labour absenteeism in Spain reached 7.6% in the first quarter of 2026, with more than 1.6 million workers estimated to be absent from their jobs on an average day. The same analysis calculates that absenteeism represented 7.2% of scheduled working hours, while temporary incapacity accounted for 5.6%.

The distinction between these numbers is important.

The 7.2% figure is hours lost, not the overall absenteeism rate. Treating the two as interchangeable would exaggerate or distort the underlying statistic.

But even with the distinction made correctly, the economic signal is difficult to ignore.

A large economy can simultaneously have people looking for work, employers looking for workers, and more than a million employed people unavailable on a given day.

That is not simply a question of unemployment. It is a question of effective labour supply.

In addition, presenteeism suggests that even this calculation remains incomplete. The worker who is absent is visible, but the worker who comes to work sick is counted as available.

Mental health has become a labour-market variable

For decades, mental health was often placed in the “wellbeing” section of corporate strategy. That category is becoming increasingly difficult to defend: the WHO estimates that approximately 15% of working-age adults were living with a mental disorder in 2019. It also identifies excessive workloads, low job control, job insecurity, discrimination, and other poor working conditions as risks to mental health.

The economic consequences extend well beyond sickness absence. A person experiencing depression or anxiety may continue working. Indeed, the WHO notes that mental-health conditions can affect productivity and the ability to retain or obtain work. This creates precisely the kind of problem that conventional HR systems are bad at seeing.

The employee is present. The employee is not necessarily dysfunctional. The employee may even appear highly committed. But concentration, memory, patience, creativity and decision quality can deteriorate. In a knowledge economy, these are not peripheral capabilities. They are the product.

The danger is greatest where errors are expensive, but output is difficult to measure. A tired surgeon, an exhausted air-traffic controller, a burned-out teacher, and an anxious financial analyst

do not produce the same kind of failure. Yet the underlying mechanism can be similar: the individual remains formally available while effective capacity deteriorates.

The WHO's conclusion is therefore economically relevant, not merely humanitarian. Good work can support mental health, while poor working conditions can undermine it. The organisation itself is part of the production environment.

The boundary between health policy and productivity policy is consequently becoming thinner.

The ageing workforce makes presenteeism more consequential

This problem would be serious even in a young labour market, but in ageing societies, it becomes structural.

The modern political response to demographic ageing has often focused on extending working lives. Retirement ages rise. Governments encourage older workers to remain economically active. Companies are urged to retain experienced employees.

All of that can make sense. But there is a missing variable. How healthy are those additional working years?

Recent European research on healthy working life expectancy shows how sharply health conditions can affect the ability to remain both healthy and employed after 50. A 2026 study covering 19 European countries found substantially lower healthy working life expectancy among people with rheumatic and musculoskeletal diseases than among those without them, with differences reaching 50% in some countries.

This matters because musculoskeletal conditions do not necessarily make work impossible. They make some forms of work much harder. A person with chronic pain may be able to work six hours but not eight. They may be able to work at a computer but not lift heavy objects. They may be productive with flexible scheduling but unable to tolerate a long commute and a rigid shift pattern.

The binary question, “Can this person work?” is therefore increasingly inadequate.

The more useful question is:

Under what conditions can this person work sustainably?

That is a very different labour-market question.

The missing middle

Most employment systems still operate around a crude binary: fit for work or not fit for work.

However, a growing part of the workforce exists somewhere between the two. They can work, but not indefinitely. They can work, but not at the previous intensity. They can work, but not in the same job design. They can work, but perhaps only four days a week. They can work, but need treatment during working hours. They can work, but not on rotating shifts. They can work, but require ergonomic changes. They can work, but need periods of recovery. They can work, but cannot absorb another year of the same workload.

This is the missing middle of labour policy – and it may contain a substantial labour reserve.

South Korea offers an instructive example from the other direction. In May 2026, 10.13 million people aged 55 to 79 were employed, the first time the number had exceeded ten million. Some 69.2% of people in that age group said they wanted to continue working in the future, with an average desired working age of 73.6.

The numbers demonstrate that ageing does not automatically mean withdrawal from work. Neither do they prove that all additional employment represents healthy, sustainable participation. Indeed, the same Korean data show that income remains a major motivation for continued work, with the average monthly pension among pension recipients in the group at 880,000 won.

That distinction matters. A society can increase the number of older people working while simultaneously failing to increase the number of older people working well. Employment statistics alone cannot tell the difference.

Presenteeism is not the opposite of commitment

There is another misconception worth challenging. The employee who comes to work while ill is not necessarily irresponsible. Quite often, it is exactly the opposite.

Presenteeism can be driven by commitment to colleagues, professional identity, financial pressure, fear of letting the team down, or an organisational culture in which absence is interpreted as weakness.

That is why simplistic attendance targets can backfire.

If managers reward attendance without understanding capacity, employees learn the wrong lesson. They learn that being physically present is more valuable than being fit to perform. The organisation may then create exactly the behaviour it later describes as a productivity problem.

The same applies to leadership. A manager who answers emails at midnight may be signalling dedication. They may also be signalling that uninterrupted availability is expected.

A team that repeatedly works through illness may look resilient. It may actually be running without a recovery mechanism.

A workplace that never loses a day to sickness may appear exceptionally healthy. It may simply have normalised working sick.

Thus, presenteeism is largely a cultural phenomenon. The question is not only whether people come to work. It is what the organisation has taught them that coming to work means.

The hidden cost is often transferred to tomorrow

One of the most important characteristics of presenteeism is that its cost may be delayed. Absence has an immediate effect. Presenteeism can have a deferred effect.

A worker comes to work sick today. They recover more slowly. The condition becomes chronic. They take longer leave later. Their performance deteriorates. They leave the organisation. The

company recruits a replacement. Knowledge disappears. The vacancy becomes difficult to fill. The remaining team absorbs the workload. The cycle begins again.

This is particularly important because research has linked sickness presenteeism with subsequent health and wellbeing outcomes, although the precise causal mechanisms vary across conditions and studies. A systematic review of longitudinal research concluded that the relationship between going to work while sick and later health outcomes deserves serious attention.

The economic lesson is straightforward. A worker who should have rested but instead works today may become a worker who cannot work tomorrow. This is why presenteeism should not be treated simply as a question of today's productivity. It is also a question of preserving tomorrow's labour supply.

The jobs themselves may be part of the problem

The conventional response to reduced capacity is often to ask whether the employee can adapt. The more productive question may be whether the job can adapt. This is especially relevant in ageing workforces.

Flexible hours can make some chronic conditions manageable. Remote or hybrid work can remove unnecessary commuting burdens. Ergonomic redesign can extend physically demanding careers. Predictable scheduling can make treatment and recovery possible. Phased returns can prevent a worker from moving directly from full absence to full workload. Task redesign can separate activities requiring physical strength from those requiring judgement and experience.

None of this means lowering standards. It means becoming more precise about what the job actually requires. The distinction is economically important because the alternative is often unnecessarily binary.

An experienced technician may not be able to work a physically demanding night shift but may be exceptionally valuable in diagnostics and training. An older nurse may not be able to sustain the same physical workload but may be able to supervise, mentor, or specialise. A teacher experiencing exhaustion may be able to remain in education under a different timetable. A worker with chronic pain may be able to contribute at high productivity for 30 hours rather than at low productivity for 40.

The question becomes one of *capacity engineering*.

From health policy to workforce intelligence

This is where the issue becomes particularly relevant for executives.

Companies already measure headcount. They measure vacancies. They measure turnover. They measure sickness absence.

Far fewer have a serious picture of effective workforce capacity.

That should change.

A modern workforce dashboard should distinguish at least four states:

-
1. Absent: the employee is not working.
 2. Present and fully capable: the employee is working at expected capacity.
 3. Present with reduced capacity: the employee is working, but health or exhaustion is affecting performance.
 4. At risk of exit: the employee remains employed but is approaching a point at which continued work in the current form is unlikely to be sustainable.
-

The last two categories are where much of the strategic information may sit.

Companies should therefore be asking different questions.

Not only:

How many people are absent?

But:

How often are people working while unwell? How many roles depend on sustained overtime? Where are error rates increasing? Where is turnover concentrated? Which occupations report the highest levels of exhaustion? Which teams have persistently high workloads? How long do employees remain on reduced capacity before leaving? How often do returns from sickness result in another absence? Which jobs could be redesigned without reducing output?

And perhaps most importantly:

Where are we counting employment when we should be counting capacity?

A national productivity problem hiding inside an HR problem

The implications extend beyond companies. Governments tend to monitor employment, unemployment, participation, and hours worked. Those indicators remain indispensable, but an ageing economy increasingly needs another measure: effective healthy working capacity. This would force policymakers to connect labour-market policy with health policy in a more serious way.

Faster access to treatment is not only a healthcare issue. Rehabilitation is not only a welfare issue. Mental-health prevention is not only a social issue. Ergonomics is not only an employer responsibility. Care infrastructure is not only a family-policy question. All of these influence the quantity and quality of labour available to the economy.

The same logic explains why raising statutory retirement ages without improving healthy working life can produce disappointing results. A government can legislate for a longer working life, but cannot legislate away chronic pain, exhaustion, depression, musculoskeletal disease, or an impossible workload.

The real objective should therefore be to increase the number of years in which people are capable and willing to work sustainably, rather than simply increasing the number of years in which they are legally permitted or economically compelled to remain attached to the labour market.

The people who are still at work

The easiest labour-market problems to see are the empty ones. The vacancy. The resignation. The unemployment statistic. The retirement. The missing worker.

Presenteeism points towards a harder category. The person who remains. The employee who keeps coming. The nurse who finishes another shift. The teacher who enters another classroom. The technician who works through pain. The manager who has stopped sleeping properly. The administrator who cannot concentrate but answers every email. The older worker who wants to continue but cannot sustain the old schedule. The employee who is technically present and economically counted, but whose capacity is steadily being consumed.

These people complicate the idea of labour supply because they are neither fully unavailable nor fully available. They occupy the space in between. And that space is becoming economically significant. The future labour crisis may therefore not be defined only by the question of **who is missing**.


It may also be defined by a more uncomfortable question:

How many of the people who show up are still healthy enough to stay?

If the answer is fewer than the employment statistics suggest, then presenteeism is not merely an HR problem.

It is a productivity problem. It is a demographic problem. It is a health-system problem. And ultimately, it is a national-capacity problem.

Five Hidden Indicators of Fatigue



1.



Sleep loss

A country may boast high participation and long operating hours while quietly running on sleep debt. That means more errors, more accidents, weaker concentration, and poorer health long before any collapse appears in the statistics.



2.



Commute strain

A long commute is often treated as an inconvenience. In reality it is a daily transfer of time and energy from households to the labor market. It shortens sleep, complicates care, and delivers workers to their jobs already depleted.



3.



Unpaid care load

Much of what keeps an economy functioning happens outside payrolls. Childcare, elder care, and domestic coordination absorb real effort. Where public services are thin, households become shock absorbers—and caregivers become chronically recovery-poor.



4.



Thermal exposure

Extreme heat rarely appears in growth statistics as a direct drag on output quality. But it slows work, raises risk, impairs judgment, and increases the bodily cost of producing the same result. Climate change is becoming an exhaustion multiplier.



5.



Interruption density

Digital systems have made many jobs more connected and more fragmented at the same time. A worker can now spend a whole day answering prompts without achieving sustained concentration. Productivity may still be recorded; attention has nonetheless been consumed.



Why it matters?



Traditional economic statistics are built to measure transactions, not depletion. They are good at counting what passed through the market and much worse at seeing what was worn down in the process. That leaves one of the most important modern scarcities under-described: not labor, exactly, but **usable human energy**.



Reading Human Capacity

Five questions every country and company should ask before assuming the workforce will be there

A country briefing usually begins with familiar questions. How fast is the economy growing? What is inflation doing? How stable is the government? What does regulation look like? Can capital move freely? Are the logistics reliable? What are the geopolitical risks?

All of these questions remain necessary. Increasingly, they are not sufficient.

For an expanding number of businesses, the more consequential question is much simpler:

Can this country actually supply the people required to make the operation work?

Not just today. Over the next three years. And not merely in terms of headcount.

Can the people be found? Can they perform the required work? Can they remain capable of doing it? Can the organisation replace them when they leave? Can the workforce adapt when technology changes, supply chains move, or geopolitical conditions deteriorate?

These questions point towards a different way of thinking about the health of an economy.

A country may have millions of working-age people and still lack technicians. It may produce large numbers of university graduates and still struggle to staff factories. It may have high employment but declining healthy working lives. It may attract foreign workers but fail to integrate them effectively. It may possess excellent workers today while quietly losing the institutions and skills required to replace them tomorrow.

The central mistake is to treat human capacity as synonymous with labour supply.

It is not.

Headcount is a stock. Human capacity is a system.

To understand that system, five questions provide a useful starting point: **availability, capability, health, continuity and adaptability.**

They should not be treated as a ranking system. They are better understood as a way of interrogating a country, a sector, or a particular investment decision.

1. Availability: Can the people actually be found?

The first question is the most obvious, but even here the conventional answer is often misleading.

A country may have a large labour force and still be difficult to staff.

The relevant question is therefore not:

How many workers exist?

It is:

How many suitable workers can be found, hired and deployed where they are needed, when they are needed?

That requires looking beyond unemployment and participation rates.

Regional labour availability matters. So does migration. So does the recognition of foreign qualifications. So do youth inactivity, gender participation gaps, commuting patterns, and the concentration of particular skills in particular cities.

A labour pool that exists statistically but cannot be mobilised operationally is not the same thing as an available workforce.

A factory may have thousands of potential workers within national borders, but almost none within a realistic commuting distance. A hospital may have plenty of qualified nurses nationally while being unable to recruit them in a particular region. A company may identify hundreds of candidates but discover that few can accept the required shifts, wage, or location.

Availability is therefore partly a question of *friction*. How much effort does it take to convert a person who exists into a person who can actually perform the job? That friction should be treated as part of the operating environment.

For a company considering an investment, this means that labour availability should be assessed at the level at which operations actually occur. Country-level unemployment is rarely enough. The relevant geography may be a city, a region or even a 45-minute commuting radius around a particular facility.

The relevant population may not be everyone of working age. It may be the number of people who possess a specific licence, qualification, language, security clearance, technical background or willingness to work particular hours.

The first question is therefore not whether labour exists. It is whether **the right labour can be mobilised**.

2. Capability: Can they actually do the work?

The second question is more uncomfortable. Suppose the workers can be found. Can they perform the work that the economy actually needs?

Educational attainment is an imperfect answer. A country can produce large numbers of graduates while employers continue to report severe shortages of applicable skills. India is an important example. The country produces millions of graduates each year, yet significant numbers of young graduates struggle to find employment compatible with their qualifications, with employers pointing to gaps between education and workplace requirements.

This is not necessarily a failure of education. It is a failure of *alignment*. The distinction matters because the future workforce will require more than accumulated credentials. Digitalisation, automation, electrification, artificial intelligence and industrial restructuring are changing the composition of demand.

A degree tells us that somebody completed a course of study. It does not necessarily tell us whether that person can troubleshoot a production system, supervise a team, work safely with advanced machinery, interpret data, solve an unfamiliar problem or learn a new process quickly.

Measures such as the OECD's PIAAC framework are therefore useful because they examine foundational adult capabilities such as literacy, numeracy and adaptive problem-solving rather than relying solely on formal educational attainment.

For businesses, capability should be considered in two dimensions:

-
1. **What can the workforce do today?**
 2. **What will it be able to do five years from now?**
-

The first is a stock, the second is a flow.

A country may have an impressive concentration of engineers today while producing too few technicians. Another may have excellent universities but weak vocational education. A third may have a large pool of technically capable workers but poor mechanisms for retraining them when an industry changes.

A workforce should therefore never be assessed only by what it knows. It should also be assessed by **how quickly it can learn what it does not yet know.**

3. Health: Can they sustain the work?

This is where conventional labour statistics become particularly weak. A worker can be employed and still have substantially reduced capacity. They can be physically present but exhausted. They can have chronic pain, depression, anxiety, or another condition that does not make employment impossible but makes conventional employment difficult. They can be working through illness, accumulating overtime, or reducing their hours simply to remain functional.

The previous article in this series examined *presenteeism* precisely because this hidden capacity loss is so difficult to see.

Thus, health needs to be treated as part of labour supply. Not as a separate social issue. Not simply as a healthcare expenditure, but as an economic variable.

The relevant question is not merely how long people live. It is how long they remain healthy enough to participate sustainably in work.

The distinction becomes increasingly important as populations age. Research on healthy working life shows that the number of years people can remain in work and good health varies substantially by occupation, socioeconomic position, and health status. Musculoskeletal disorders, for example, can significantly shorten healthy working lives even when they do not make employment impossible.

This changes the meaning of a labour shortage. Imagine two countries with identical numbers of workers aged 55 to 64. In the first, workers have good access to healthcare, flexible employment, effective rehabilitation and jobs that can be adapted to changing physical capacity. In the second, treatment is slow, work is rigid and physically demanding occupations are difficult to modify. The demographic statistics are identical. The available workforce is not.

Health also has a company-level dimension. An organisation should know not only how many employees it has, but where capacity is being consumed. Persistent sickness absence, repeated short-term leave, injuries, burnout, overtime dependence, and unusually high turnover may indicate that an organisation is operating beyond the sustainable capacity of its workforce.

This is where the concept of presenteeism becomes particularly useful. The employee who is absent is visible. The employee who is present but functioning at reduced capacity is much harder to see. A country with low unemployment and high employment may therefore still have a considerable amount of unused human capacity. The question is whether institutions can recover it.

4. Continuity: Can the workforce replace itself?

The fourth question concerns time. Even a highly capable workforce eventually retires, migrates, changes profession, or dies.

The strategic question is whether the system can replace what it loses.

This is more complicated than calculating the number of people entering and leaving employment. Organisations lose tacit knowledge when experienced people leave. Industries lose specialised skills that may take years to reproduce. Public institutions lose institutional memory. Companies discover that a job description can be replaced on paper while the person who actually knew how the system worked cannot.

Continuity therefore asks:

Can knowledge, leadership, and specialised capability reproduce themselves?

Can a retiring engineer be replaced? Can a senior nurse train enough successors? Can a technical supervisor be promoted from within? Can an organisation replace its local managers without importing an entirely new management structure? Can an industrial ecosystem preserve knowledge when an entire generation approaches retirement?

Demography makes these questions increasingly urgent. For business, these are not simply demographic statistics. They are **replacement risks**.

A country may remain wealthy, technologically sophisticated, and politically stable while becoming increasingly dependent on a shrinking number of experienced people. That creates a different kind of vulnerability.

The critical question becomes not:

How many people work here?

But:

What happens when the people who know how this system works are no longer here?

5. Adaptability: Can the workforce absorb disruption?

The final question may be the most important.

No workforce remains adequate forever. Technology changes. Industries move. Energy systems transform. Supply chains are reconfigured. Geopolitical relationships deteriorate. Artificial intelligence changes tasks that once seemed permanent.

The resilient workforce is therefore not necessarily the one with the highest current skill level. It is the one capable of **absorbing change without prolonged loss of productive capacity**.

This is adaptability.

It includes vocational education, lifelong learning, employer-sponsored training, occupational mobility, digital competence, and the ability to move workers from declining activities into adjacent ones.

The World Economic Forum's *Future of Jobs Report 2025* identified technological change, geoeconomic fragmentation, demographic shifts, and the green transition among the principal forces reshaping employment through 2030. Their significance is easy to underestimate.

If an economy loses 10,000 jobs in one sector and creates 10,000 new jobs elsewhere, the transition is not necessarily economically neutral.

The people who lost their jobs may not possess the skills required for the new ones. The regions may be different. The wages may be different. The training may take years. The result can be simultaneous unemployment and labour shortage.

That is why adaptability should be viewed as a form of economic infrastructure. A country with excellent universities but weak retraining may be less resilient than it appears. A country with strong vocational institutions, apprenticeships, and employer-linked training may be considerably better positioned than conventional human-capital statistics suggest.

The critical question is therefore:

How quickly can the workforce change without breaking?

AI changes all five questions

Artificial intelligence should not be treated as a separate category alongside these five dimensions. In fact, it cuts through all of them.

AI can reduce the number of people required for some tasks. It can increase the productivity of others. It can eliminate routine entry-level work while increasing demand for people capable of supervising, integrating, and applying advanced systems. It can preserve institutional knowledge by making information easier to retrieve. It can also make some workers obsolete if neither the individual nor the organisation can adapt. This creates a particularly important distinction between ***substitution and augmentation***.

The question is not simply:

How many jobs will AI destroy?

It is:

Where will AI increase the effective capacity of human beings, and where will it make existing human capabilities less valuable?

The answer will differ by occupation and country.

A technically skilled worker using AI-assisted diagnostics may become more productive. A junior administrative worker performing routine information processing may face greater substitution pressure. A doctor may use AI to handle parts of diagnostic work while becoming more important in judgement, communication and responsibility. A manager may automate reporting while discovering that leadership, trust and coordination have become more important rather than less.

AI therefore interacts with every dimension. It changes who needs to be available. It changes what capability means. It changes cognitive workload and potentially health. It changes how institutional knowledge can be preserved. And it raises the premium on adaptability.

The countries that benefit most may not be those with the largest populations or even the highest levels of current education. They may be those best able to **combine human capability with machines**.

From five questions to one picture

These five dimensions become most useful when considered together. A country can be strong in one and weak in another.

It can have...

**high availability, but poor capability;
strong capability, but declining health;
excellent workers today, but weak continuity;
large human numbers, but low deployability;
strong domestic talent, but limited ability to attract outsiders;
high technological exposure, but weak adaptability.**

This is why a single ranking is often less useful than a structured diagnosis.

Consider a hypothetical investment. Country A has a large young population, high graduate numbers, and low labour costs. However, its vocational system is weak and employers report severe difficulties finding experienced technicians. Country B has a much smaller workforce, higher wages and an older population. Yet it has excellent vocational training, high labour productivity, strong health infrastructure, and sophisticated automation.

Which country has the stronger human-capacity position? There is no answer without knowing what the business needs. For a labour-intensive garment factory, the calculation may favour A. For an automated semiconductor facility requiring specialist technicians, B may be considerably stronger.

The analytical task is therefore to connect **human capacity with the actual operating model**.

That is where a serious *external intelligence* approach becomes useful. The relevant information is scattered across labour statistics, demographic projections, health data, migration figures, education systems, employer surveys, regional data, wage movements, industrial policy, and technology trends.

The challenge is not merely collecting it – it is connecting it.

The practical test

For executives evaluating a country, sector, or investment, five questions provide a useful starting framework:



1 Availability

Can we recruit the people we need within the required geography and timeframe?



2 Capability

Do they possess the skills required today, and can the local system produce the skills required tomorrow?



3 Health

Can this workforce sustain the workload, physically and mentally, for the duration our business model assumes?



4 Continuity

What happens when experienced workers retire, leave or become unavailable? Where is the succession pipeline?



5 Adaptability

How quickly can this workforce absorb AI, automation, regulatory change, industrial restructuring or geopolitical disruption?

The answers should then be **tested against the specific business model.**



A company requiring **500 warehouse workers** should not analyse human capacity in the same way as a company requiring **50 semiconductor technicians**.



A hospital, data centre, defence contractor and software company do not consume the same workforce.



Neither do they face the same risks.



This is precisely why broad labour-market indicators should be treated as **starting signals rather than conclusions.**

The deeper analysis begins when the headline statistic **stops being reassuring.**





THE VACANCY THAT REFUSES TO DIE

A mini case study in vacancy persistence

Imagine a mid-sized industrial company trying to hire one controls technician for a regional production site.

The role is first posted in January 2025. The salary is benchmarked to the company's standard technical pay band. The requirements include five years' experience, familiarity with three proprietary systems, weekend availability, and a 45-minute commuting radius. HR receives applicants, but few are technically credible. Hiring managers reject most of them.

By April, the role has been reposted twice. Recruiters widen the search area. Two candidates reach final interviews; one declines on compensation, another accepts a competitor's offer. Overtime for the existing maintenance team rises. Planned preventive maintenance begins slipping.

By August, the vacancy has now been "open" in some form for eight months, but in public data it may appear as multiple job ads, making demand look stronger than it is. Internally, the company still talks about a "single unfilled role". Operationally, however, that one vacancy is already multiplying into late repairs, contractor costs, management attention, and higher burnout among incumbents.

At the one-year mark, the company has changed the title, loosened one credential, added a sign-on bonus and started using an external staffing firm. Still no hire. The site manager now delays a planned capacity upgrade because there is no confidence the plant can support additional automation without technical labour to maintain it.

At 18 months, leadership finally redesigns the role. Instead of searching for a fully formed technician, the company splits the work into three layers: a senior diagnostics specialist covering multiple sites, a locally based technician trained on core systems, and remote vendor support for software escalation. It also raises pay, funds training, and changes shift structure.

What should an executive learn from this story?

First, the vacancy total was misleading. There was not a stream of healthy labour demand. There was one persistent mismatch.

Second, the critical metric was not “how many vacancies do we have?” but “how long has this capability gap been impairing operations?”

Third, the vacancy itself became a strategic forcing mechanism. Left unresolved, it changed capex timing, maintenance reliability, workforce burnout and the company’s operating model.

The vacancy refused to die because the original job design no longer matched the labour market reality.

Why time-to-fill is often more informative than vacancy totals?

If vacancy totals describe the stock of unfilled jobs, time-to-fill measures the duration of non-resolution. For many executive purposes, duration is the more informative metric.

A position that takes 20 days to fill in a high-turnover market is not equivalent to one that takes 240 days to fill in a low-turnover market. The second has much larger implications for productivity, service delivery, expansion planning, and managerial bandwidth. Academic work on vacancy duration has shown that **job vacancy duration carries information about future firm performance**, precisely because long-duration vacancies reveal something about organisational frictions, market conditions, or role quality that raw counts do not.

Time-to-fill also helps separate demand strength from hiring dysfunction. If vacancy counts remain high but time-to-fill falls, the market may be functioning reasonably well despite elevated demand. If vacancy counts are moderate but time-to-fill rises sharply, the labour market is probably becoming less efficient at matching. For decision-makers, that distinction is essential.



The Missing Apprentice

What if the first thing AI automates is the process by which humans become experts?

There is a familiar way to tell the AI story, and it is flattering to executives. The machine arrives. The drudgery disappears. A junior analyst produces in ten minutes what used to take half a day. A service team clears its queue faster. A bank drafts its compliance report sooner. A software team ships more with fewer people. The dashboard improves. Costs ease. Productivity rises. The case appears self-evident.

And in many organisations, at least for now, it is.

But there is another story unfolding beneath the visible gains. It is quieter, slower, and far more consequential. It concerns the roles that are no longer refilled, the graduate schemes that are trimmed, the junior rungs that are thinned out because AI can now perform enough of the work once given to beginners. The short-term result is attractive: lower costs and higher output. The delayed result may be much harder to repair: fewer experienced professionals five, ten, or fifteen years from now.

That is the real tension. AI can raise productivity while simultaneously eroding the capability base on which future productivity depends. In other words, the first thing AI may automate is not expertise itself, but the process through which humans acquire it. That is the missing apprentice problem.

The job may vanish before the profession does

Senior professionals are not disappearing en masse. In fact, many are becoming more valuable. Organisations still need judgment, political reading, client trust, accountability and contextual interpretation. Those are precisely the capabilities that current AI systems do not reliably replicate. But the work beneath those senior roles — the repetitive, preparatory, first-pass work through which younger employees used to learn the craft — is now under direct pressure.

This matters because most professions do not reproduce themselves by magic. They reproduce themselves through a ladder. A newcomer starts with bounded tasks. Then comes repetition. Then correction. Then pattern recognition. Then exceptions. Then responsibility. Then, eventually, judgment.

Take away enough of the bounded work at the bottom, and the ladder still looks intact from the top — for a while. But eventually, there are fewer people climbing it.

This is no longer just a philosophical concern. Recent reporting drawing on Stanford Digital Economy Lab findings says the employment shortfall for U.S. workers aged 22 to 25 in highly AI-exposed occupations widened from 15% in mid-2025 to about 19% by June 2026, with reduced hiring accounting for most of the shift rather than a wave of firings. That distinction matters. The door is not slamming on experienced workers first. It is simply opening less often for new ones.

Other markets are showing a similar pattern. In South Korea, the central bank reported that of 285,000 youth jobs lost over four years, 94% were in AI-exposed sectors. In the same broad category of sectors, workers aged 50 and above gained 230,000 jobs. The implication is striking: younger workers are losing the entry ramp while more experienced workers remain relatively protected.

In India's SaaS sector, revenues have continued to rise while headcount has barely moved. The top 20 high-growth SaaS unicorns reportedly grew headcount by less than 1% over 12 months to March 2026, even as the market crossed \$13 billion in revenue. Entry-level openings fell, junior technology positions declined 35% between FY25 and FY26 across major markets, and recruiters describe a new pattern in which backfills are delayed for three to six months, with many roles never reopening at all. One recruiter's line is especially telling: "The most telling signal we see is not layoffs. It is the pause before a backfill."

That is exactly what a capability problem looks like in its earliest phase. Not a dramatic collapse. A missing replacement.

Productivity is real, but so is the hollowing-out risk

None of this means the productivity case for AI is false. On the contrary, it is often very real.

A well-known study of 5,179 customer-support agents found that access to a generative AI assistant increased productivity by roughly 14%, with the biggest gains accruing to novice and lower-skilled workers. In some summaries of the same work, the least skilled improved by around 35% on certain measures. This is one of the most important results in the whole debate, because it captures both the promise and the danger of AI at once. AI can help less experienced workers perform more like

stronger ones. That is excellent news for output now. It is more complicated news for the organisational logic of employing large junior cohorts later.

Banking offers the same dual story. In Europe, generative AI is estimated to be able to add between €180 billion and €306 billion in value globally, equivalent to roughly 9% to 15% of current operating profits in the sector. Intesa Sanpaolo has reportedly cut the time required for EU-wide stress-test reports by more than 80%, while its customer-service chatbot resolves around 85% of text-based requests. Yet the same programme is linked to 9,000 job cuts by 2027, around 10% of staff, offset by 3,500 new digitally skilled hires.

This is the central arithmetic of the moment. AI does not merely eliminate labour in the abstract. It often compresses the middle and lower layers of the professional pyramid while preserving, and in some cases increasing, demand for people who can supervise, interpret, validate, and decide.

The legal profession is beginning to describe the outcome in precisely these terms. One recent discussion of law firms argues that AI attacks the “grinder” layer — the research, drafting and analysis work done by juniors — thereby creating a “hollow pyramid” and a “verification gap.” If the junior no longer produces enough material directly, how exactly do they learn to recognise weak reasoning, missing facts, or strategic error?

This is why the issue is larger than employment statistics alone. It goes directly to the reproduction of professional judgment.

The old inefficiency was actually a training system

Many firms have long treated junior work as an economic imperfection: necessary, but inelegant. New graduates were slow, supervision was expensive, and the real value seemed to sit higher up the chain.

But that view misses something important. The apprentice model was never just a staffing arrangement. It was a capability-production system.

Entry-level roles did not exist merely because firms needed someone to tidy slides, conduct preliminary research, prepare first drafts, reconcile data, or compile briefing packs. They existed because those tasks, done repeatedly and corrected by more experienced colleagues, formed the cognitive grammar of a profession. The repetition mattered. The corrections mattered. The exposure to edge cases mattered.

What AI is unusually good at automating are precisely those repetitive and semi-routine tasks that once justified employing beginners in the first place.

In consulting, policy, law, software, intelligence analysis, banking, compliance, communications, and administration, the formative tasks have often been the most automatable ones: retrieval, synthesis, comparison, drafting, formatting, summarising, coding standard patterns, checking documents, preparing background notes. These tasks were not glamorous. They were developmental.

If those tasks disappear without deliberate redesign of the training path, the firm may preserve throughput while degrading formation. That is why the problem is best framed not as job loss, but as transmission failure.

The most dangerous capability risks rarely announce themselves clearly. They appear first in ordinary managerial decisions that each seem rational on their own.

Do not refill the departing analyst; the team now has copilots. Cut the graduate intake this year; productivity is up. Slow the promotion round; AI has flattened the middle. Concentrate review in a smaller group of top performers; machine output still needs experienced eyes. Each decision makes local sense. Collectively, they thin the ladder.

Augmentation can still erode development

One of the most common rebuttals is that AI does not replace young workers; it augments them. Sometimes that is true, but it depends entirely on how the augmentation is designed.

If AI helps a junior understand better, compare faster, receive richer feedback, and see more examples under supervision, then it can accelerate development. If AI simply allows them to produce acceptable-looking outputs without deeply understanding the underlying reasoning, then productivity rises while capability formation weakens.

This should be a strategic lesson for firms. Automation of tasks and automation of development are not the same thing. A business may rationally automate first drafts. It is much riskier to automate away the sequence of supervised repetitions through which someone learns what a good draft is, why it is good, and how it fails under pressure.

If the junior employee's job becomes prompting, packaging, and passing along machine outputs they could not independently recreate or challenge, they may look highly effective while remaining thinly developed. The organisation then mistakes fluency for mastery.

A nation can consume its future talent stock

Because this newsletter issue is *Health of Nations*, it is worth widening the lens. The missing apprentice is not merely a company problem. It is a national-capacity problem.

Countries do not remain competitive because they possess a static stock of talented people. They remain competitive because they continuously convert the inexperienced into the experienced. That requires ladders: in engineering, software, medicine, public administration, intelligence, manufacturing, finance, logistics, teaching, and law.

If AI narrows too many of those ladders at the entry point, advanced economies may discover — too late — that they preserved margins while damaging the process through which future competence is produced.

This risk is especially important for societies already dealing with demographic ageing and strategic competition. A country with a thinner pipeline of developed professionals becomes more dependent on a shrinking stock of incumbents, more exposed to external talent shortages, and more vulnerable to institutional brittleness.

A business can save money by consuming its future. A country can do the same.

The surprising perspective

Most public debate still asks whether AI will replace experts. That may be the wrong question. A more surprising and more strategically important question is whether AI will replace the conditions under which experts are made before it replaces experts themselves.

That would explain the current pattern rather well. Senior people remain valuable. Experienced workers remain needed. AI improves output. Firms become more selective. Graduate hiring weakens. Entry-level work is redesigned upward or removed altogether. The visible gains arrive quickly. The capability erosion is deferred.

By the time the consequences are obvious, the organisation has already spent years living off a stock of previously trained people. That is what makes the missing apprentice problem so treacherous. The dashboard can improve while the institution quietly weakens.

What healthy organisations should do?

The answer is not to preserve pointless drudge work for sentimental reasons. It is to preserve the developmental function that drudge work once served. That requires deliberate redesign.

If AI can do the routine work, juniors need structured exposure to verification, exception-handling, supervised judgment, and client or stakeholder context earlier than before. One useful practical idea drawn from recent commentary is the “exception rotation”: first, junior staff should verify AI-generated work against source materials, then review edge cases with more experienced colleagues. Finally, they would take responsibility for bounded categories of exceptions. In other words, the repetitions change, but they do not disappear.

Firms should also stop treating non-refilled junior roles, reduced graduate intakes, and flatter promotion funnels as mere HR details. They are leading indicators of capability erosion. In sectors where tacit knowledge matters, boards should ask a blunt question: are we improving productivity, or merely liquidating part of our future talent pipeline?

The healthiest response will be neither anti-AI nostalgia nor blind techno-optimism. It will be capability-preserving adoption: using AI to raise human output without destroying the institutional ladders through which judgment is formed.

The real test of an AI-native firm

A genuinely strong AI-native organisation will not be one that simply uses fewer juniors. It will be one that can answer a harder question: how, in an AI-rich environment, do we still produce the next generation of trusted professionals? That is the test.

Because if the first thing AI automates is the process by which humans become experts, then the cost will not show up immediately in layoffs statistics or quarterly earnings. It will show up later — when the firm cannot find enough strong managers, trusted analysts, senior operators, credible advisers, or experienced decision-makers who actually learned the craft from the ground up.

At that point, the savings will already have been booked – and the apprentice will already be missing.

Scenario Piece
Four News Dispatches from the Future
2035



The Automated Nation, 2035



The morning clinics are calm, not because demand has fallen, but because software has absorbed the invisible work that once consumed half the day. Triage is continuous. Records update themselves. Home sensors flag risks before a fall becomes an admission. In factories, ports, warehouses, and municipal services, fewer workers supervise more machines. The workforce is smaller and older, but less overwhelmed than expected. Yet in the quiet rooms of dementia wards and special-needs classrooms, one truth remains: the last mile is still human. The debate is no longer whether AI saved the nation from demographic decline. It is whether the nation used AI to protect human beings, or merely to stretch them further.

This is the most attractive quadrant for governments that want to avoid political conflict over migration while sustaining productivity. But in reality, very few countries can occupy it fully, because automation works best in administration, logistics, manufacturing, diagnostics support, and some routine services; it works much less cleanly in eldercare, childrearing, teaching, and relational healthcare.

The Migrant Nation, 2035



The airport is now part labour market, part civic frontier. Fast-track visas bring in nurses, coders, electricians, eldercare aides, and graduate students who will become all four in sequence. Population decline has stopped in the capital, but not in the provinces. New languages fill public schools; foreign-born clinicians keep hospitals open; political coalitions reorganize around the meaning of belonging. Success and strain arrive together. In prosperous districts, the model looks elegant: talent in, growth up, ageing moderated. In overloaded housing markets and underplanned suburbs, it looks fragile. The country has solved the arithmetic of labour, but not yet the grammar of solidarity.

This quadrant is economically powerful but politically difficult. It requires housing supply, school capacity, healthcare access, credential recognition, local integration, and a civic narrative of belonging. When those conditions hold, migration can refresh age structures and keep essential sectors functioning. When they do not, the same model becomes politically brittle.

The Protected Nation, 2035



The state speaks a new language now. Maternity wards, semiconductors, geriatric prevention, housing design, school meals, and industrial strategy appear in the same budget speech. Family policy is no longer sentimental; it is strategic. The care economy is treated as productive infrastructure. Employers redesign jobs for 60-year-olds; cities retrofit neighbourhoods so 80-year-olds can remain independent; health ministries are judged not only by hospital waiting times but by years of healthy life retained. Birth rates have not dramatically rebounded. That was never the whole point. The point was to stop governing as if human capacity were an abundant resource.

It does not mean avoiding migration or automation. It means prioritizing and preserving domestic human capacity through family policy, healthy longevity, care reform, and labour-market redesign—and do so in a way that the public broadly understands as fair, practical, and future-oriented.

The Exhausted Nation, 2035



The spreadsheets look reassuring. Output per worker has improved. Export sectors remain globally competitive. Markets are calm. But schools have merged, clinics have thinned, parents have delayed children again, and eldercare is carried by daughters, migrants, and fatigue. More citizens live alone; more regions feel old before their time. The nation has not collapsed. It has adapted, technically. But adaptation has taken the form of silent subtraction: fewer siblings, fewer carers, fewer births, fewer public guarantees, fewer assumptions that tomorrow will be easier than today. The political question arrives late, after the economic one seemed settled: when does a society become too tired to reproduce itself confidently?

This is where a country appears to be functioning, even performing adequately on GDP or productivity, but is doing so by drawing down social reserves. It is the land of delayed births, exhausted careers, older workers in worsening health, migrants filling gaps without full incorporation, and widening regional abandonment.

Subscribe to the Mission Grey Guild Newsletter

<https://missiongrey.com/insights/#newsletters>

Guild Members



1

Richard Andrew

Senior APAC public affairs and geopolitical strategy expert

Geopolitics

2

Runi Irnandari Arnowo

Senior banking executive and former Managing Director at Deutsche Bank, with 30+ years across corporate banking, transaction banking and capital markets in Asia-Pacific

Finance Strategy Operations

3

Marino Auffant

Harvard-awarded historian and strategist

Geopolitics Strategy

4

Stephen J Blank

Senior foreign policy expert and fellow

Geopolitics

5

Mathieu Boulègue

Transatlantic security researcher and consultant

Geopolitics Security

6

Gloria Chen

Multidimensional geopolitical analyst with a focus on space geopolitics

Geopolitics

7

Rori Comiskey

US Special Forces officer with operational intelligence expertise

Geopolitics Security Operations

8

Kelly Couto

Latin American specialist and business consultant

Geopolitics Finance Operations

9

Zsuzsa Anna Ferenczy

EU-China and Indo-Pacific specialist, academic and former European Parliament political advisor

Geopolitics Regulation Strategy

10

Fabio Frettoli

International security analyst and project designer

Geopolitics Security

11

Joseph Gozoya

PhD candidate in economics, focusing on tax data, policy, and labour markets

Finance Regulation

12

Samantha Karlin

Global empowerment strategist and CEO

Geopolitics Strategy

13

Alan Leung

Enterprise risk and geopolitical intelligence advisor

Geopolitics Finance Security Strategy

14

Looi Teck Kheong

Global legal and political strategist, ASEAN specialist

Geopolitics Regulation

15

Ludvig Lund

Seasoned security consultant

Security

16

Dorota Maczuga

East Asia expert and co-founder of Polylocal

Geopolitics

17

Andrew Marshall

Global media, IP and commercial strategy leader

Strategy Operations Finance

18

Shabab Muhaddes

FinTech and financial risk executive

Finance Technology

19

Toshihisa Nagai

Geoeconomic risk and regulatory intelligence strategist focused on sanctions, export controls, and capital architecture across APAC and the EU

Geopolitics Regulation Finance Technology

20

Fernando Prats

LATAM political, regulatory and public affairs advisor

Geopolitics Regulation Strategy

21

Rwizi Rweizooba Ainomugisha

Freelance wargame researcher, analyst and designer

Geopolitics Strategy Technology

22

Amandine Schneickert

International strategic policy executive and lobbyist

Regulation Geopolitics

23

Stéphane Schwander

Independent sports business advisor; former FEI and UEFA executive

Strategy Finance

24

Scott R. Siler

Political risk advisor and strategy professor

Geopolitics Regulation Strategy

25

Diana The Hui Ling

Senior legal, insurance, and risk management expert

Finance Regulation

26

Daniel Valerio

Former EY Partner in charge of Americas Retail Risk and Audit

Finance Operations

27

Pekka Väisänen

Political scientist focused on French and European affairs

Geopolitics Regulation

28

William Vogt

Emerging Markets Expert, with an emphasis on Sino-Lusophone (China/Portuguese-speaking) Affairs.

Technology Geopolitics

29

Pekka Virkki

Guildmaster, geopolitical author and journalist

Geopolitics

